

Neural Relational Inference for Interacting Systems

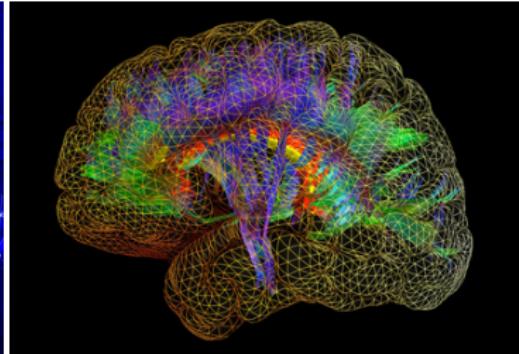
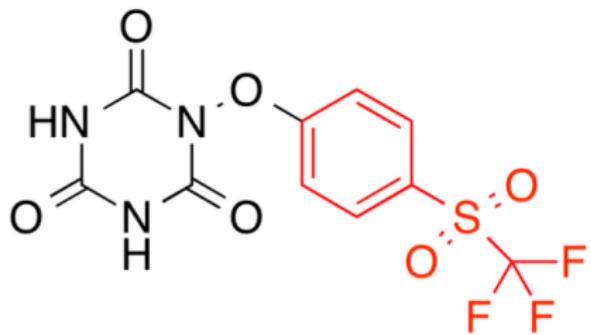
Petar Veličković

Artificial Intelligence Group
Department of Computer Science and Technology, University of Cambridge, UK

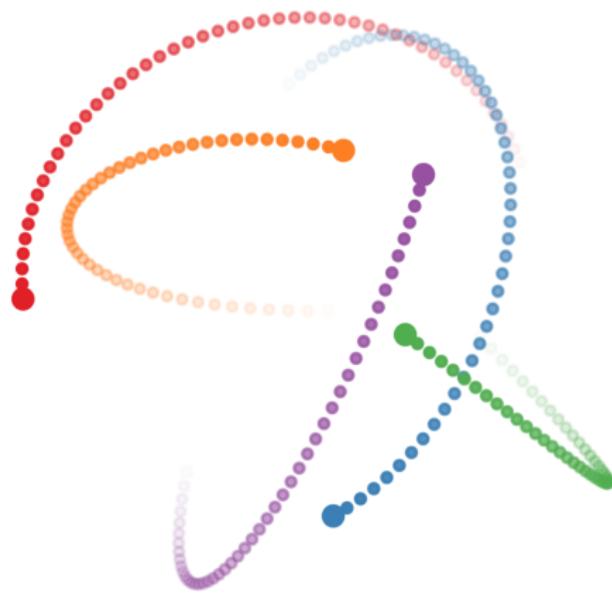
Introduction

- ▶ In this talk, I will survey the recently published **Neural Relational Inference** model (Kipf, Fetaya *et al.*, ICML 2018).
- ▶ This model enables the discovery and exploitation of latent interactions between objects, through the synergy of *graph convolutional networks* and *variational autoencoders*.
- ▶ Exciting results + avenues for further work!

Graphs are **everywhere**!



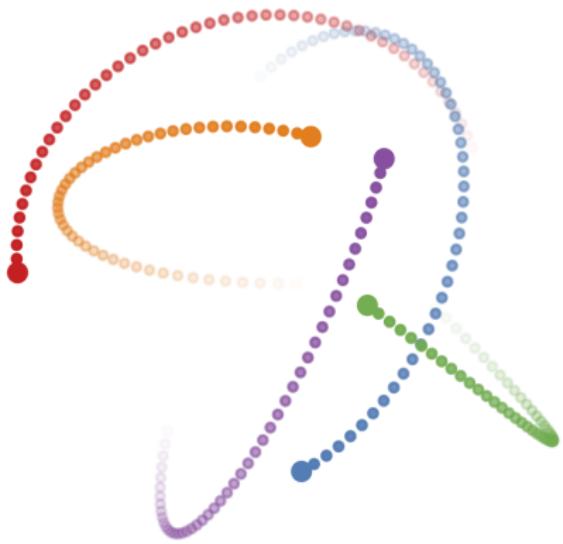
...but can we always *see* them?



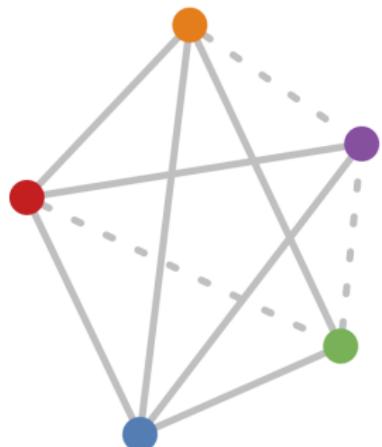
Relational inference

- ▶ Virtually all graph convolutional techniques require a graph to be *provided* as input!
- ▶ However, often we will only have access to *node features*...
- ▶ Approaches such as *Relational Networks* (Santoro *et al.*, 2017), or *VAIN* (Hoshen, 2017) circumvent this by assuming a **complete graph** (i.e. all-pairs interactions).
- ▶ But most interaction graphs have properties (such as *sparsity*) that we may wish to explicitly demand!
- ▶ Furthermore, we may wish to *identify* and *decouple* different **types of interaction**.

Our task for today



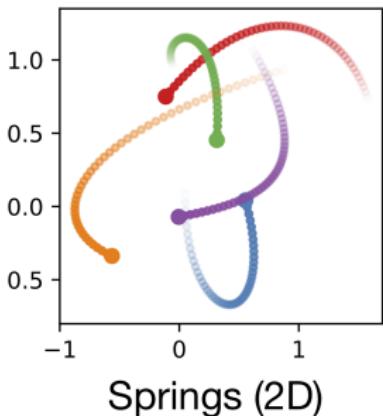
Observed dynamics



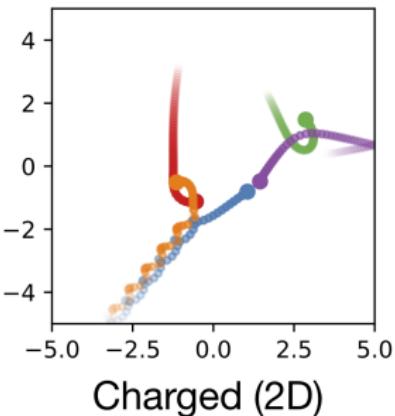
Interaction graph

Motivation: predicting *trajectories*

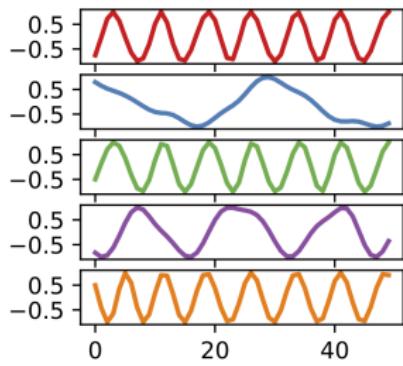
- ▶ **Input:** Trajectories (e.g. coordinates) $\vec{x}_i^{\leq t}$ for each particle i .
- ▶ **Output:** Future trajectories $\vec{x}_i^{>t}$ for each particle i .



Spring (2D)



Charged (2D)



Kuramoto (1D)

- ▶ The interaction graph between particles will be a *byproduct*!

Simple baseline #1: RNN

- ▶ Let \vec{x}^t denote the coordinates of all particles at time t :

$$\vec{x}^t = [\vec{x}_1^t, \vec{x}_2^t, \dots, \vec{x}_n^t]$$

- ▶ We can now define a recurrent neural network (e.g. LSTM or GRU) to operate on this sequential input:

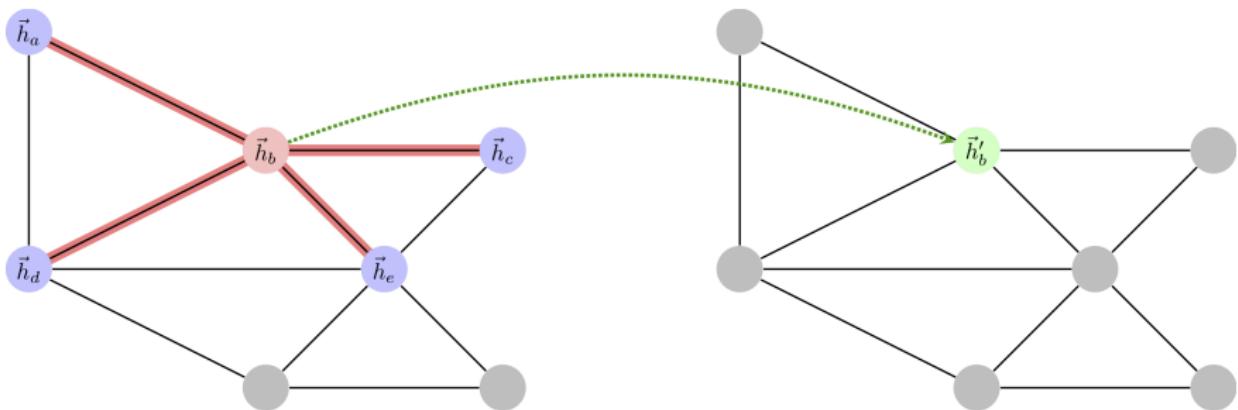
$$\vec{h}^t = RNN(\vec{h}^{t-1}, \vec{x}^t)$$

- ▶ From its hidden states, we can predict the future timesteps:

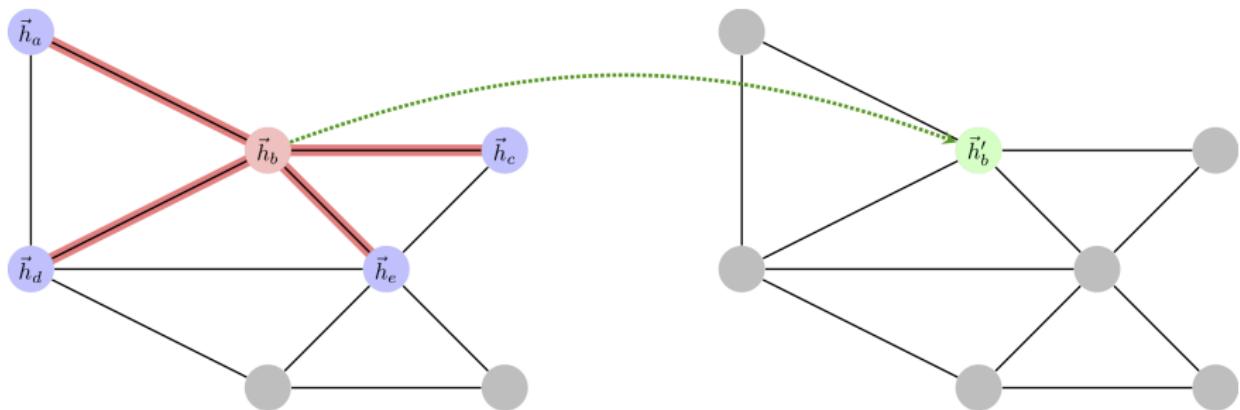
$$\vec{x}^{t+1} = f(\vec{h}^{t+1})$$

where f is an MLP.

Graph convolutional network



Graph convolutional network



In a nutshell, obtain higher-level representations of a node i by leveraging its *neighbourhood*, \mathcal{N}_i !

$$\vec{h}_i^{\ell+1} = g^\ell(\vec{h}_a^\ell, \vec{h}_b^\ell, \vec{h}_c^\ell, \dots) \quad (a, b, c, \dots \in \mathcal{N}_i)$$

where g^ℓ is the ℓ -th *graph convolutional layer*.

The MPNN framework

- ▶ The NRI model leverages a graph convolutional layer inspired by *message-passing neural networks* (Gilmer *et al.*, 2017).

The MPNN framework

- ▶ The NRI model leverages a graph convolutional layer inspired by *message-passing neural networks* (Gilmer *et al.*, 2017).
- ▶ First, compute *edge messages*, $\vec{h}_{i \rightarrow j}^\ell$, for each edge $i \rightarrow j$ in the graph. Apply a simple MLP, f_e^ℓ , over the features of i and j :

$$\vec{h}_{i \rightarrow j}^\ell = f_e^\ell(\vec{h}_i^\ell, \vec{h}_j^\ell)$$

The MPNN framework

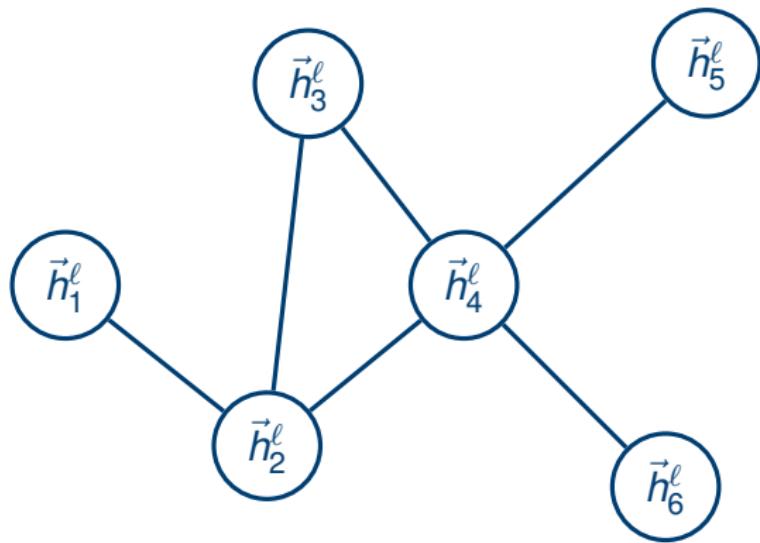
- ▶ The NRI model leverages a graph convolutional layer inspired by *message-passing neural networks* (Gilmer *et al.*, 2017).
- ▶ First, compute *edge messages*, $\vec{h}_{i \rightarrow j}^\ell$, for each edge $i \rightarrow j$ in the graph. Apply a simple MLP, f_e^ℓ , over the features of i and j :

$$\vec{h}_{i \rightarrow j}^\ell = f_e^\ell(\vec{h}_i^\ell, \vec{h}_j^\ell)$$

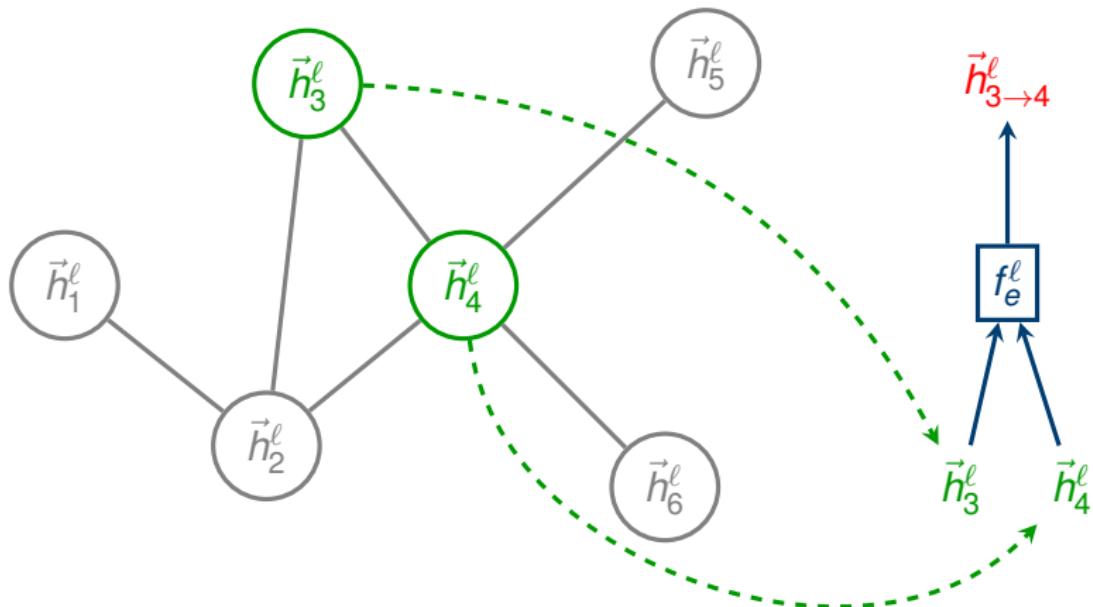
- ▶ Then, aggregate all messages entering a node j to obtain the next-level features, $\vec{h}_j^{\ell+1}$. Apply a simple MLP, f_v^ℓ , over the summed messages.

$$\vec{h}_j^{\ell+1} = f_v^\ell \left(\sum_{j \in \mathcal{N}_i} \vec{h}_{i \rightarrow j}^\ell \right)$$

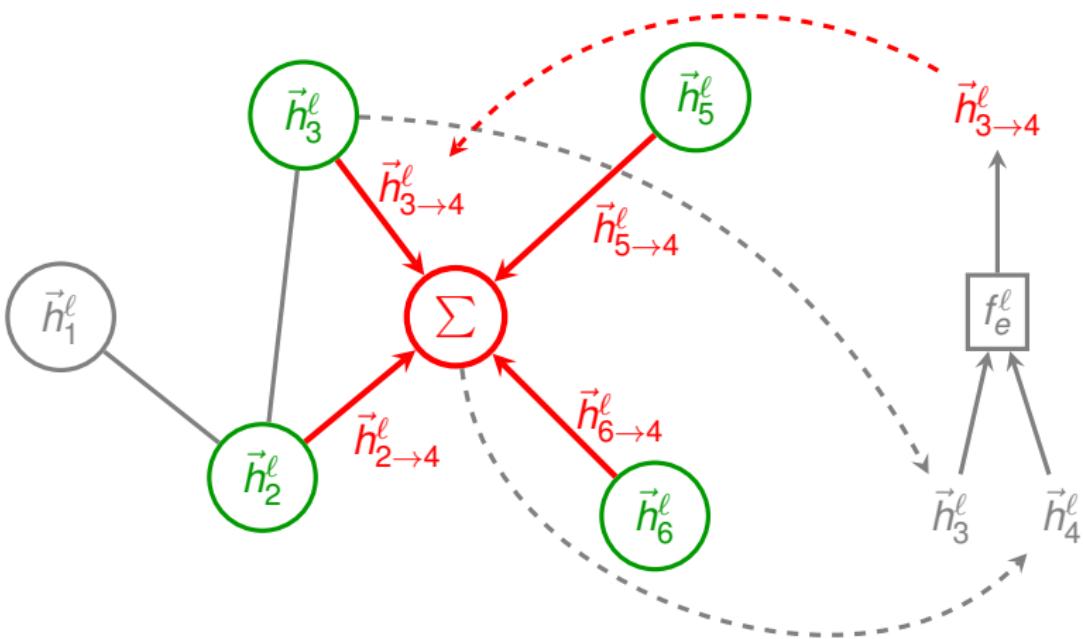
MPNN: initial setup



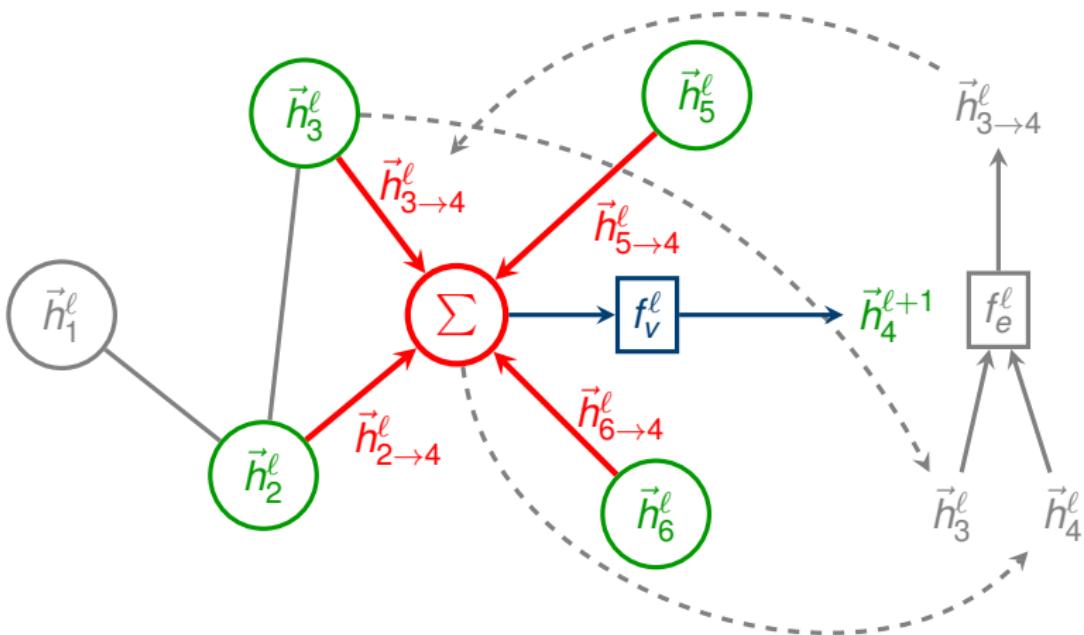
MPNN, computing messages



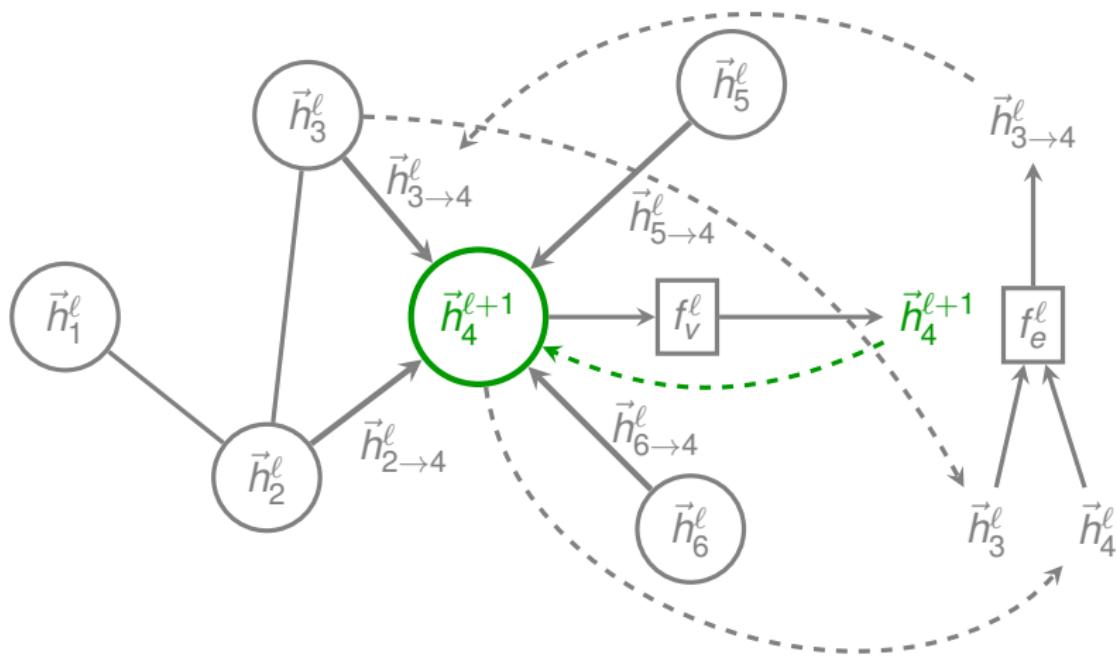
MPNN, aggregating messages



MPNN, computing node features



MPNN, next-level features



Simple baseline #2: Complete graph

- ▶ As another baseline approach, we may use this kind of layer to predict trajectories using a complete graph (assume all pairs of nodes interact).
- ▶ The equations of the baseline become equivalent to:

$$\begin{aligned}\vec{h}_{i \rightarrow j}^t &= f_e(\vec{x}_i^t, \vec{x}_j^t) \\ \vec{x}_j^{t+1} &= f_v \left(\sum_{i \neq j} \vec{h}_{i \rightarrow j}^t \right)\end{aligned}$$

with a few kinks, specific to the trajectory predicting task...

Simple baseline #2: Complete graph

- ▶ First, to simplify the job of the network, have it only predict *changes in position*:

$$\vec{h}_{i \rightarrow j}^t = f_e(\vec{x}_i^t, \vec{x}_j^t)$$

$$\vec{x}_j^{t+1} = \boxed{\vec{x}_j^t} + f_v \left(\sum_{i \neq j} \vec{h}_{i \rightarrow j}^t \right)$$

Simple baseline #2: Complete graph

- Also, *explicitly model uncertainty*; will be useful for the variational framework later on.

$$\vec{h}_{i \rightarrow j}^t = f_e(\vec{x}_i^t, \vec{x}_j^t)$$

$$\boxed{\vec{\mu}_j^{t+1}} = \vec{x}_j^t + f_v \left(\sum_{i \neq j} \vec{h}_{i \rightarrow j}^t \right)$$

$$\boxed{\vec{x}_j^{t+1} \sim \mathcal{N}(\vec{\mu}_j^{t+1}, \sigma^2 \mathbf{I})}$$

Simple baseline #2: Complete graph, *with GRU*

- The model thus far assumed the *Markov property* (i.e. that \vec{x}^{t+1} depends fully on \vec{x}^t). This is OK for physics, but if necessary, we can alleviate the constraint by using a recurrent update:

$$\vec{h}_{i \rightarrow j}^t = f_e(\vec{x}_i^t, \vec{x}_j^t)$$

$$\boxed{\vec{h}_j^{t+1} = GRU \left(\left[\vec{x}_j^t, \sum_{i \neq j} \vec{h}_{i \rightarrow j}^t \right], \vec{h}_j^t \right)}$$

$$\vec{\mu}_j^{t+1} = \vec{x}_j^t + f_v \left(\vec{h}_j^{t+1} \right)$$

$$\vec{x}_j^{t+1} \sim \mathcal{N}(\vec{\mu}_j^{t+1}, \sigma^2 \mathbf{I})$$

Interaction graph

- ▶ This baseline can be improved if we specify an explicit *interaction graph*. Initially, assume there are K edge types (with one type reserved for “no edge”).
- ▶ Then, define a binary tensor $\mathbf{z} \in \mathbb{R}^{V \times V \times K}$ such that z_{ijk} denotes whether the edge $i \rightarrow j$ is of the k -th type.
- ▶ Assume an edge cannot have more than one type, i.e., \vec{z}_{ij} is one-hot.

Leveraging the interaction graph

- Now this graph can be exploited—define a separate MLP f_e^k for each edge type. For the Markov decoder:

$$\begin{aligned}\vec{h}_{i \rightarrow j}^t &= f_e(\vec{x}_i^t, \vec{x}_j^t) \\ \vec{\mu}_j^{t+1} &= \vec{x}_j^t + f_v \left(\sum_{i \neq j} \vec{h}_{i \rightarrow j}^t \right) \\ \vec{x}_j^{t+1} &\sim \mathcal{N}(\vec{\mu}_j^{t+1}, \sigma^2 \mathbf{I})\end{aligned}$$

The NRI *decoder*

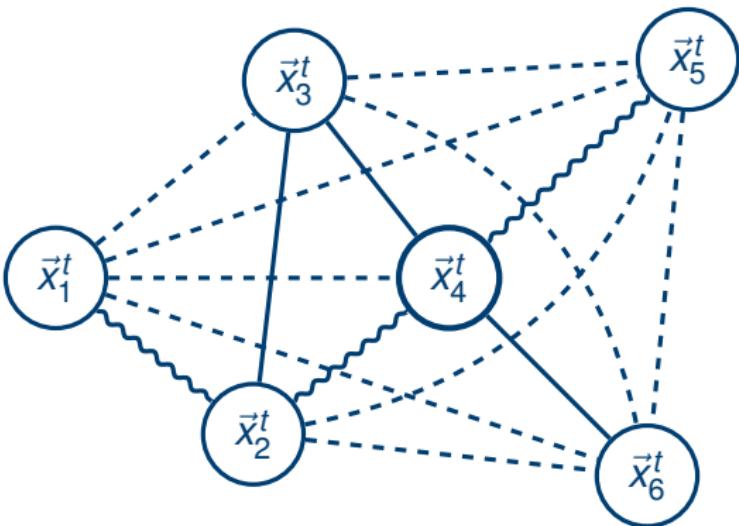
- Now this graph can be exploited—define a separate MLP f_e^k for each edge type. For the Markov decoder:

$$\vec{h}_{i \rightarrow j}^t = \boxed{\sum_k z_{ijk} f_e^k(\vec{x}_i^t, \vec{x}_j^t)}$$

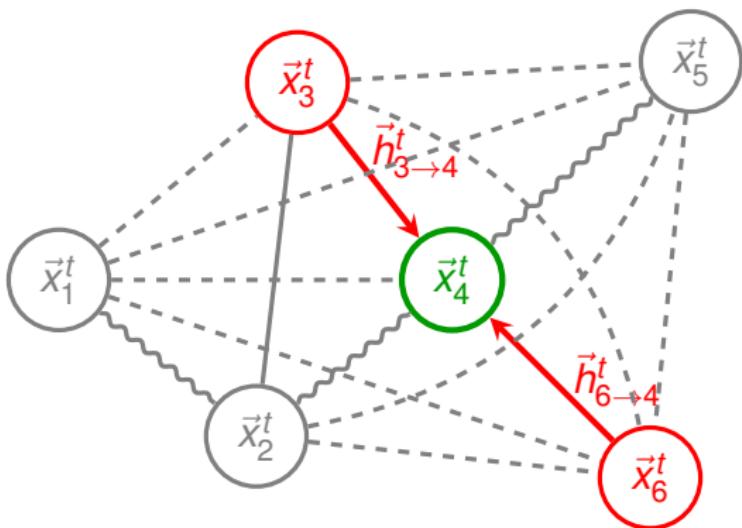
$$\vec{\mu}_j^{t+1} = \vec{x}_j^t + f_v \left(\sum_{i \neq j} \vec{h}_{i \rightarrow j}^t \right)$$

$$\vec{x}_j^{t+1} \sim \mathcal{N}(\vec{\mu}_j^{t+1}, \sigma^2 \mathbf{I})$$

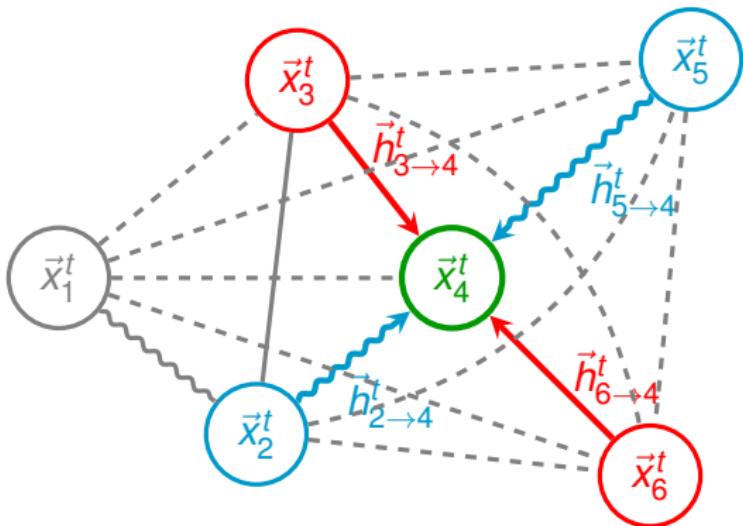
The NRI *decoder*



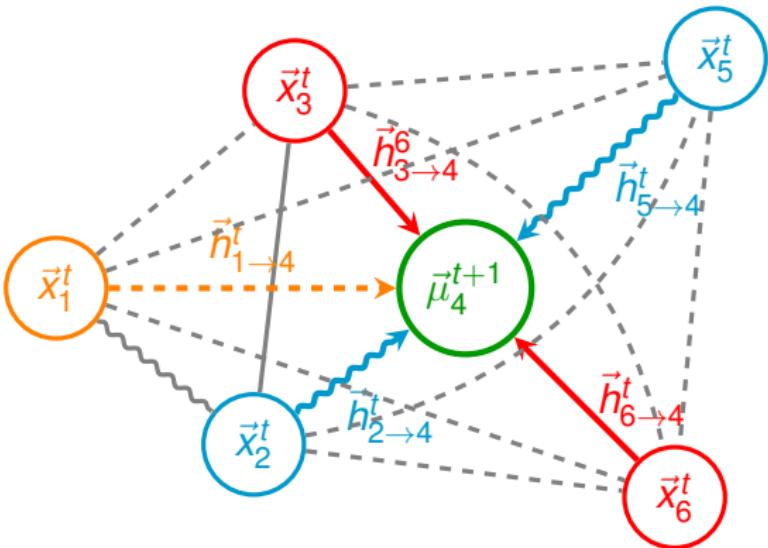
The NRI *decoder*, computing messages...



The NRI *decoder*, computing messages...



The NRI *decoder*, computing messages...



Latent graph inference

- ▶ We are still tasked with discovering the entries of the tensor \mathbf{z} .
- ▶ **Idea:** Use MPNNs over a complete graph once more—then classify edge types based on the edge messages $\vec{h}_{i \rightarrow j}$.
- ▶ This time, *stack two layers*—so that edges can be derived based on global interactions!
 - ▶ $\vec{h}_{i \rightarrow j}^1$ will only depend on \vec{x}_i and \vec{x}_j ;
 - ▶ $\vec{h}_{i \rightarrow j}^2$ will depend on all the nodes in the graph.

The NRI *encoder*

- In equation form:

$$\vec{h}_j^1 = f(\vec{x}_j)$$

$$\vec{h}_{i \rightarrow j}^1 = f_e^1(\vec{h}_i^1, \vec{h}_j^1)$$

$$\vec{h}_j^2 = f_v^1 \left(\sum_{i \neq j} \vec{h}_{i \rightarrow j}^1 \right)$$

$$\vec{h}_{i \rightarrow j}^2 = f_e^2(\vec{h}_i^2, \vec{h}_j^2)$$

$$z_{ij} \sim \text{Categorical}(\text{softmax}(\vec{h}_{i \rightarrow j}^2))$$

where f is an embedding, and f_e^1 , f_v^1 and f_e^2 are MLPs.

The variational setup

- ▶ The encoder gives us the probability distribution $q(\mathbf{z}|\vec{x})$, and the decoder gives us the probability distribution $p(\vec{x}|\mathbf{z})$.
- ▶ Combine learning the two in a VAE-style framework by maximising the evidence lower bound (ELBO):

$$\mathcal{L} = \underbrace{\mathbb{E}_{\mathbf{z} \sim q(\mathbf{z}|\vec{x})} [\log p(\vec{x}|\mathbf{z})]}_{\text{Reconstruction accuracy}} - \underbrace{D_{KL}(q(\mathbf{z}|\vec{x}) \| p(\mathbf{z}))}_{\text{Regularisation}}$$

- ▶ The prior $p(\mathbf{z})$ can encode desirable properties of the latent graph. **Sparsity** is enforced by setting the probability of “no edge” to be higher than the other types.

Backpropagating through the sampling

- ▶ The operation of selecting z_{ij} is a *discrete* decision—therefore, we cannot directly propagate gradients through it.
- ▶ Can use the **Gumbel softmax** trick to circumvent this:

$$\vec{z}_{ij} = \text{softmax}((\vec{h}_{i \rightarrow j}^2 + \vec{g})/\tau)$$

where $g_k \sim \text{Gumbel}(0, 1)$ and τ is a temperature parameter (converges to one-hot when $\tau \rightarrow 0$).

- ▶ This is a *continuous approximation* to the discrete distribution—and gradients can be propagated through it.

Avoiding degenerate decoders

- ▶ Optimising the ELBO directly would involve only *single-step* predictions (predicting \vec{x}^{t+1} from \vec{x}^t). This can often be nicely approximated by ignoring relational structure altogether!
- ▶ To enforce robust decoders, predict many steps at once! Every M steps, feed back the ground-truth input.

Avoiding degenerate decoders, *cont'd*

$$\vec{\mu}_j^2 = \text{decode}(\vec{x}_j^1)$$

$$\vec{\mu}_j^3 = \text{decode}(\vec{\mu}_j^2)$$

$$\vec{\mu}_j^4 = \text{decode}(\vec{\mu}_j^3)$$

⋮

$$\vec{\mu}_j^{M+1} = \text{decode}(\vec{\mu}_j^M)$$

$$\vec{\mu}_j^{M+2} = \text{decode}(\vec{x}_j^{M+1})$$

$$\vec{\mu}_j^{M+3} = \text{decode}(\vec{\mu}_j^{M+2})$$

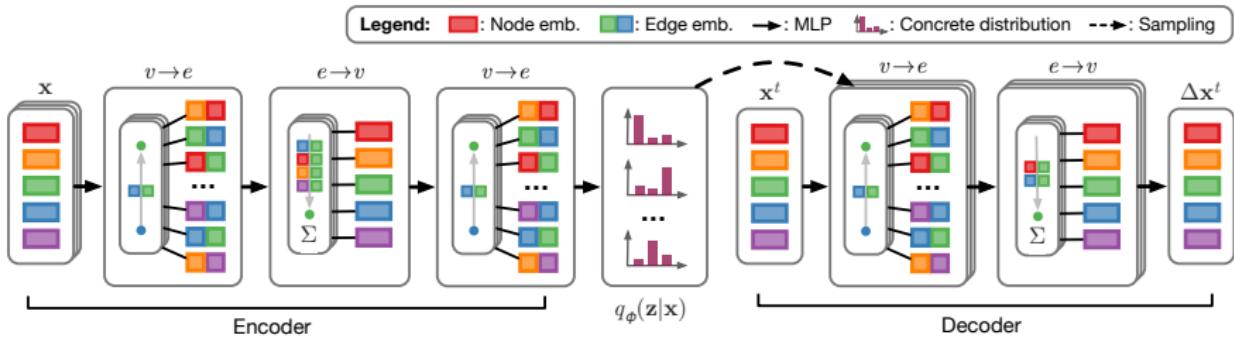
⋮

Putting it all together

For a given training trajectory, \vec{x} , of length T :

1. Compute $q(\mathbf{z}|\vec{x})$ using the encoder.
2. Sample \vec{z}_{ij} from $q(\mathbf{z}|\vec{x})$ using the Gumbel softmax trick.
3. Execute the decoder to obtain $\vec{\mu}^t$ for $t \in \{2, 3, \dots, T\}$.
4. Compute the reconstruction error (of $\vec{\mu}^t$ against \vec{x}^t) and KL-divergence (of $q(\mathbf{z}|\vec{x})$ against the prior $p(\mathbf{z})$).
5. Optimise the ELBO using gradient descent.

The NRI architecture



Physics simulations: latent graph discovery

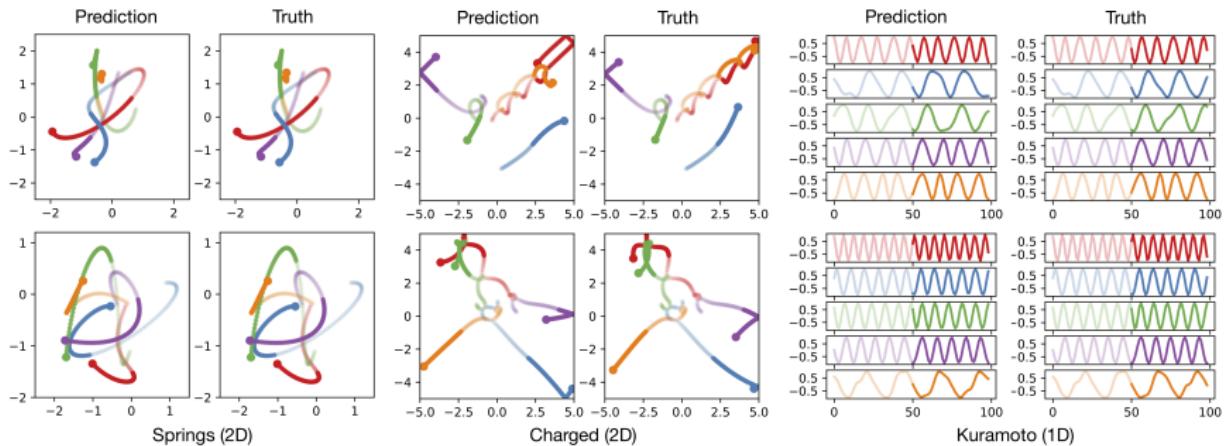
| Model | Springs | Charged | Kuramoto |
|---------------|-----------------------|-----------------------|-----------------------|
| 5 objects | | | |
| Corr. (path) | 52.4 \pm 0.0 | 55.8 \pm 0.0 | 62.8 \pm 0.0 |
| Corr. (LSTM) | 52.7 \pm 0.9 | 54.2 \pm 2.0 | 54.4 \pm 0.5 |
| NRI (sim.) | 99.8 \pm 0.0 | 59.6 \pm 0.8 | — |
| NRI (learned) | 99.9 \pm 0.0 | 82.1 \pm 0.6 | 96.0 \pm 0.1 |
| Supervised | 99.9 \pm 0.0 | 95.0 \pm 0.3 | 99.7 \pm 0.0 |
| 10 objects | | | |
| Corr. (path) | 50.4 \pm 0.0 | 51.4 \pm 0.0 | 59.3 \pm 0.0 |
| Corr. (LSTM) | 54.9 \pm 1.0 | 52.7 \pm 0.2 | 56.2 \pm 0.7 |
| NRI (sim.) | 98.2 \pm 0.0 | 53.7 \pm 0.8 | — |
| NRI (learned) | 98.4 \pm 0.0 | 70.8 \pm 0.4 | 75.7 \pm 0.3 |
| Supervised | 98.8 \pm 0.0 | 94.6 \pm 0.2 | 97.1 \pm 0.1 |

Physics simulations: trajectory prediction

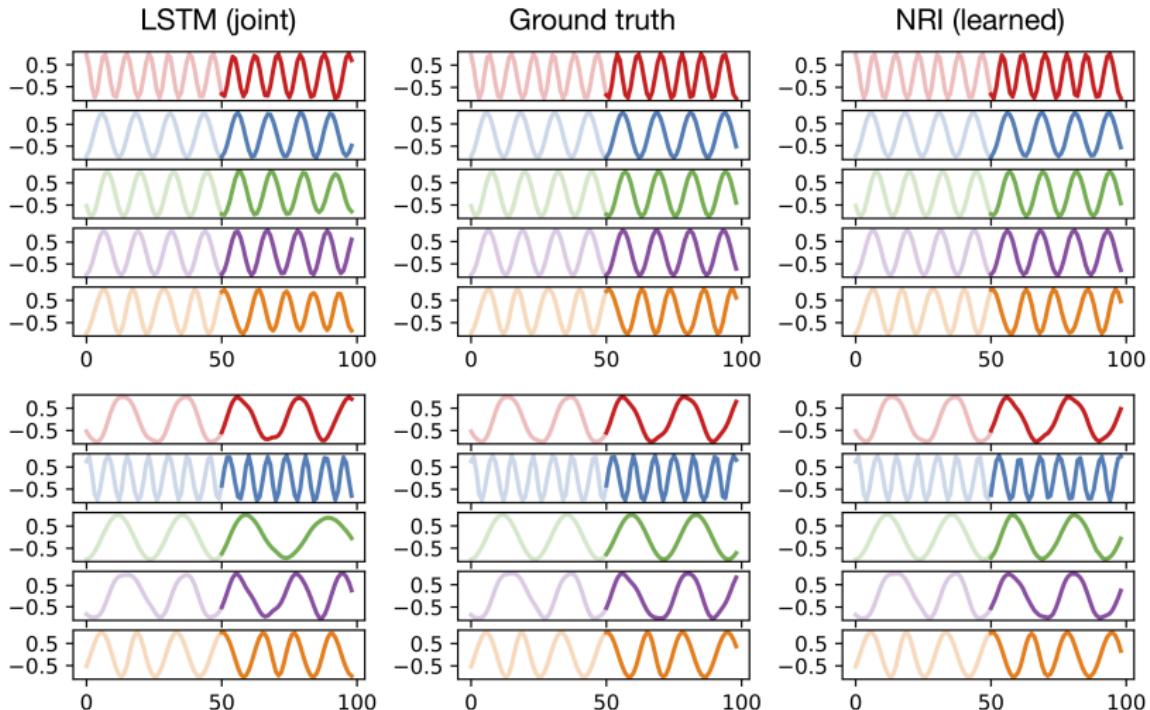
| Prediction steps | Springs | | | Charged | | | Kuramoto | | |
|------------------|----------------|----------------|----------------|----------------|----------------|----------------|----------------|----------------|----------------|
| | 1 | 10 | 20 | 1 | 10 | 20 | 1 | 10 | 20 |
| Static | 7.93e-5 | 7.59e-3 | 2.82e-2 | 5.09e-3 | 2.26e-2 | 5.42e-2 | 5.75e-2 | 3.79e-1 | 3.39e-1 |
| LSTM (single) | 2.27e-6 | 4.69e-4 | 4.90e-3 | 2.71e-3 | 7.05e-3 | 1.65e-2 | 7.81e-4 | 3.80e-2 | 8.08e-2 |
| LSTM (joint) | 4.13e-8 | 2.19e-5 | 7.02e-4 | 1.68e-3 | 6.45e-3 | 1.49e-2 | 3.44e-4 | 1.29e-2 | 4.74e-2 |
| NRI (full graph) | 1.66e-5 | 1.64e-3 | 6.31e-3 | 1.09e-3 | 3.78e-3 | 9.24e-3 | 2.15e-2 | 5.19e-2 | 8.96e-2 |
| NRI (learned) | 3.12e-8 | 3.29e-6 | 2.13e-5 | 1.05e-3 | 3.21e-3 | 7.06e-3 | 1.40e-2 | 2.01e-2 | 3.26e-2 |
| NRI (true graph) | 1.69e-11 | 1.32e-9 | 7.06e-6 | 1.04e-3 | 3.03e-3 | 5.71e-3 | 1.35e-2 | 1.54e-2 | 2.19e-2 |

It might seem as if the LSTM outperforms the NRI on Kuramoto!
Qualitative analysis may show otherwise...

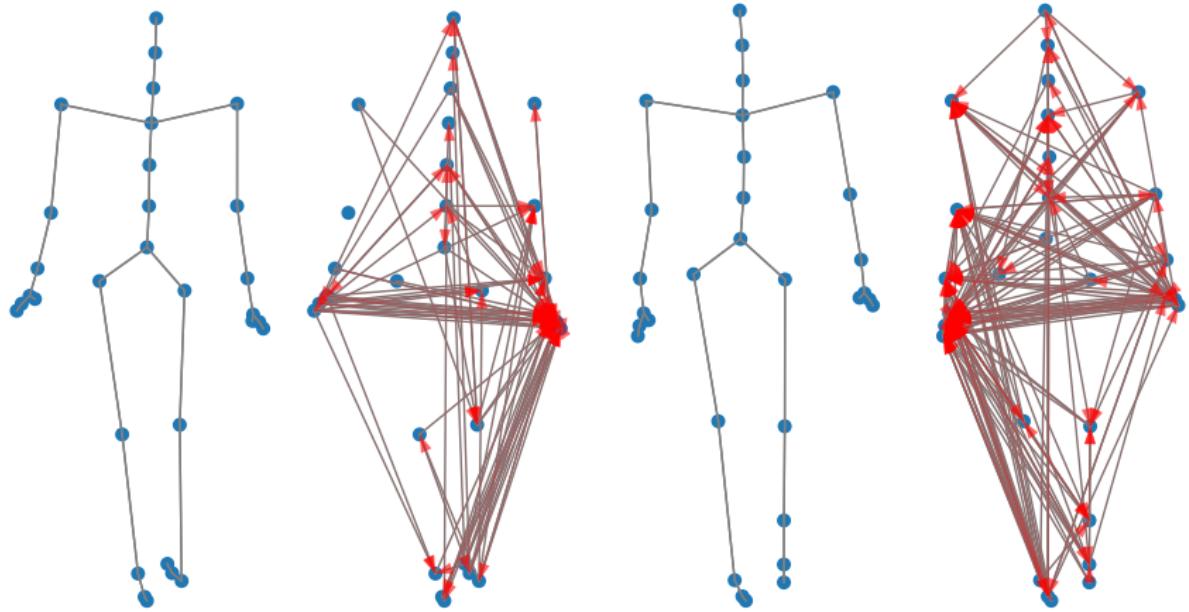
Physics simulations: qualitative results



Physics simulations: qualitative results

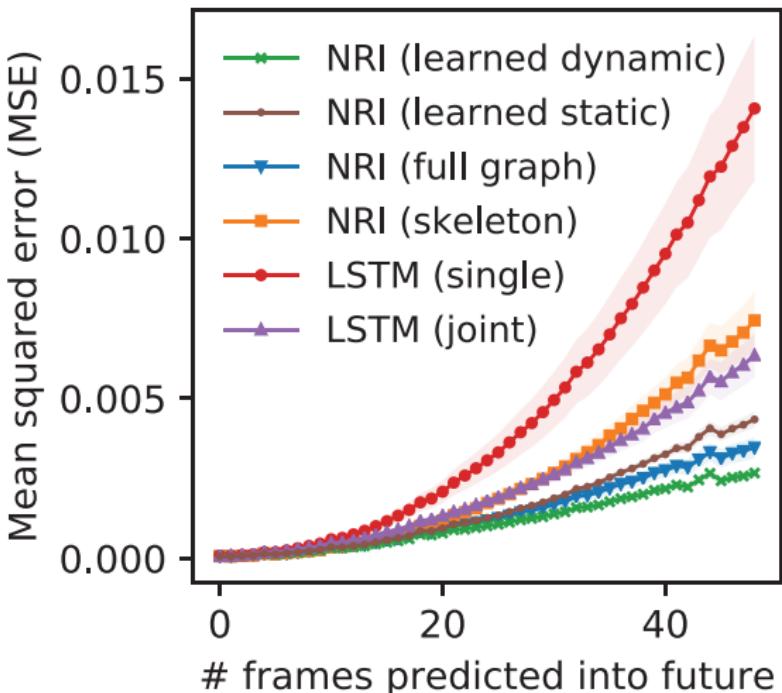


Motion capture

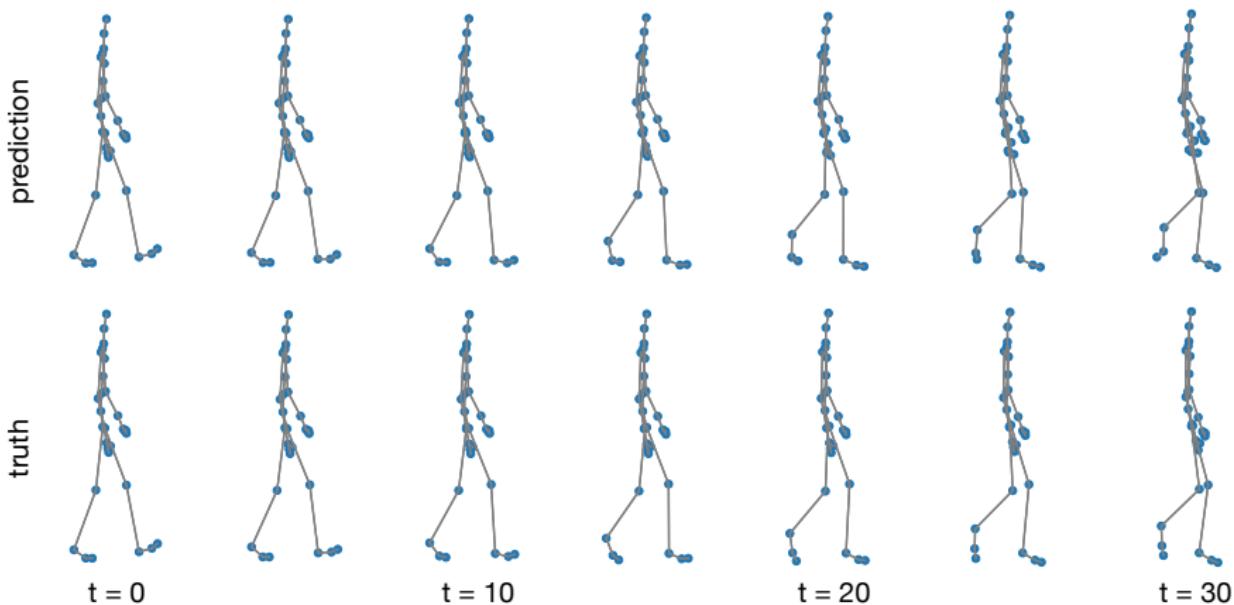


The graph is now **dynamic!** Re-evaluate at every decoding step.

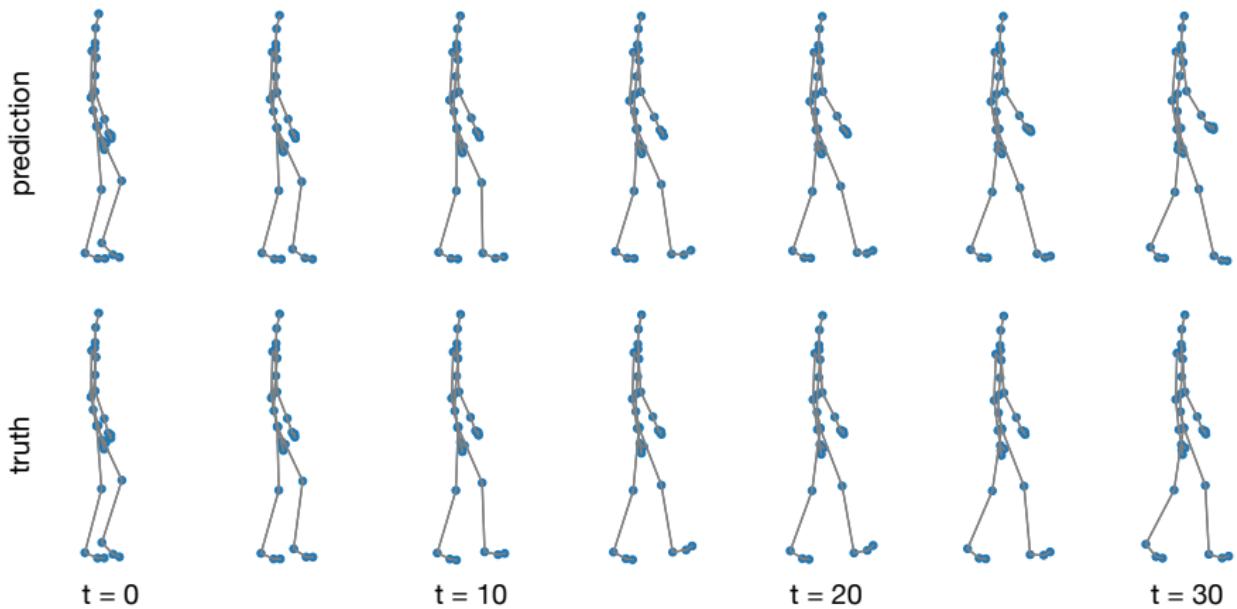
Motion capture: trajectory prediction



Motion capture: qualitative results



Motion capture: qualitative results



Concluding remarks

- ▶ The NRI is an extremely versatile model for inferring latent interaction graphs from pointwise trajectories.
- ▶ Latent graph discovery is still in its early phases of development—plentiful improvements possible!
- ▶ **Limitation:** *does not scale to large graphs!* $O(V^2)$ memory requirements, and computing edge messages makes subsampling cumbersome.
- ▶ Should not be required—most real-world graphs are sparse! But techniques we have thus far need to start with complete graph, and gradually discover sparsity...

Thank you!

Questions?

petar.velickovic@cst.cam.ac.uk

<http://www.cst.cam.ac.uk/~pv273/>